

CS 228 Problem Set 1

Hugh Zhang

February 19, 2017

Problem 1

Let the structure be a two node binary Bayes net with one edge from A to B. Let $P(A = 1) = 0.51$, $P(B = 1 | A = 1) = .5$, $P(B = 1 | A = 0) = .99$

Then, $P(B = 1) = 0.51 * .5 + 0.49 * .99 = .7401$, so the most likely outcome based on marginals without looking at conditionals is $A = 1, B = 1$. However, the $P(A = 1, B = 1) = .51 * .5 = .255$. This is not as likely as the most likely outcome, $A=0, B=1$ which has $P(A = 0, B = 1) = .49 * .99 = .4851$.

Problem 2

Basic message passing is nd^2 for a chain. Labeling the nodes from A to Y WLOG, you can factor

$$P(A \dots Y) = \phi(A, B) * \phi(B, C) \dots \phi(X, Y)$$

Z, which is the normalizing constant is just

$$\sum_A \dots \sum_Y P(A \dots Y) = \sum_A \sum_B \phi(A, B) * \sum_C \phi(B, C) \dots \sum_Y \phi(X, Y)$$

Z takes nd^2 time to calculate because this is just message passing across the entire array. Note that when we try calculating the marginals for a pair of variables, the variables cut out three parts of the chain. WLOG call them x_i and x_j with $i < j$. For all x_k with $k < i$ or x_m with $m > j$, then you can do basic message propagation to calculate each chain for nd^2 using basic message passing. Then, to get the middle part [bottom of the equation listed below], you can choose to extend the left side (WLOG) and calculate the middle chain for each fixed value of x_i , and then unite it with the right chain, then normalize to have a complete probability. Thus, since you do message passing d times when you try to marginalize the variable out (once for each value of the variable), it takes nd^3 the middle calculation still dominates the run time. To calculate all n^2 pairs, the runtime is then n^3d^3

Then

$$P(x) = \frac{l(x)\phi(X_i, x_{i+1}) \dots \phi(x_{j-1}, X_j) * r(x)}{l(x) \sum_{x_i} \phi(x_i, x_{i+1}) \dots \sum_{x_j} \phi(x_j - 1, x_j) * r(x)}$$

2.2

Note that X_i and X_j are independent conditioned on X_{j-1}

$$\sum_{x_{j-1}} P(X_i, X_{j-1}) P(X_j | X_{j-1}) = \sum_{x_{j-1}} P(X_i, X_{j-1}, X_j) = P(X_i, X_j)$$

2.3

I claim you can do this in $O(n^2 d^3)$ time. First you calculate Z , which takes $O(nd^2)$, which is fine. Then we calculate all the individual marginals $P(x_i)$ for all i . This is also just basic message passing, so its $O(nd^2)$. Then, you calculate all adjacent pairs using our above algorithm. $P(x_1, x_2) \dots P(x_{n-1}, x_n)$. This takes nd^3 time per pair, and N pairs, so we are at $O(n^2 d^3)$. Then, to calculate all pairs of distance x apart given that you've calculated all pairs of distance $x-1$ apart and below, we use the formula above. (This is the inductive step)

$$\begin{aligned} P(X_i, X_j) &= \sum_{x_{j-1}} P(X_i, X_{j-1}) P(X_j | X_{j-1}) \\ &= \sum_{x_{j-1}} P(X_i, X_{j-1}) \frac{P(X_j, X_{j-1})}{P(X_{j-1})} \end{aligned}$$

We have all of these values, since $j - i - 1 = x - 1 < x$, and we can use our previously calculated values. There are d^2 possible values of X_i and X_j , and N possible pairs of distance X . Marginalizing over all values of X_{j-1} adds another factor of d to our calculations for a total of nd^3 per inductive step. Since we have n inductive steps, this is a total of $n^2 d^3$, exactly what we want.

Problem 3

If you change a factor (by changing its value, or by adding an edge) in a clique, check all its neighbors and see if they are affected, AKA the separation set contains the factor that has been changed. If they are affected, then everything in that entire subtree needs to be updated, since they all depend on that factor in the message passing. If the factor is not in the sep-set, then it gets marginalized out anyways.

3.2

Pick any clique that contains X . Once you find the probability of all the variables in the clique, you can marginalize out the other variables so that you have $P(x)$. The markov blanket for this, is all the neighboring cliques. If the factor you changed is not inside this markov blanket, you don't need to do anything since it all gets marginalized out anyways.

If your factor is inside your markov blanket, then you just need to make sure that all the incoming messages to the clique that contains X is correct, since $P(C)$ is just the factor times the product of all the incoming messages with marginalization. If your factor is inside the clique, you do not need to update any messages and can just directly compute the new marginal with your new factor. This is because all the messages coming in are correct (they do not depend on any messages going out of this clique or they would double count). If you are in among the neighbors, just updating the message inward (if the factor is in the sep-set) is sufficient.

Problem 4

4.2.1

$P(X \mid E)$ is computationally intractable (NP hard to calculate on a general bayes net). Sampling from this distribution is also hard, since you can't calculate $P(X \mid E)$ for any given X easily.

Computing $P(X, E)$ though is doable, since you just walk through the entire Bayes tree and propagate down on your probabilities. Forwardprop

b

1

Notice that belief propagation runs exponential in the size of the largest clique, but since we have a handy dandy tree structure, it runs in polynomial time. This suggests a way for us to sample well.

Transform your Bayesian net to a Markov net. Since every node has exactly one parent, there are no V structures so the transition is natural. Then, we can run variable elimination to get $P(X_i \mid E)$ for all X_i . From here, we need a way to sample everything instead of just finding the MAP assignment. But we can do this in a clever way. Topologically sort the variables you need, then sample the first one. Then, ELIMINATE IT from the graph and rerun Variable Elimination to compute the new marginals on the new graph, or essentially $P(x_j \mid E, x_i)$. Since there are N variables to eliminate, and every BP operation is polynomial, then this is polynomial.

So:

For each variable X (topologically sorted):

- Rerun Belief Propagation on the new graph
- Sample from X 's possible values according to X 's marginal probability on this new graph
- Eliminate the variable using Variable Elimination to create the next graph.

At the end, you will have $P(X_1 | E) * P(X_2 | E, X_1) \dots$ and a valid sample

2

$$w(x) = \frac{P(X = x, E = e)}{Q(X = x | E = e)}$$

$W(x)$ is supposed to approximate $P(e)$. In the lecture notes, the Q function models $P(X | E)$, but here it models P , so we divide instead by $Q(X = x | E = e)$

You can calculate $P(X = x, E = e)$ easily through forward prop, and $Q(X = x | E = e)$ through BP.

3

$$\hat{P}(X = x | E = e) = \frac{\sum_{t=1}^T \delta_{x_i} w(z^t)}{\sum_{t=1}^T w(z^t)}$$

where δ is your indicator function to see if your assignment is consistent with your evidence, and w is the function defined in part 2. The idea is the top approximates $P(X, E)$ and the bottom estimates $P(E)$, and thus importance sampling samples from $P(X | E)$

Problem 5

a

CODE HERE XX

```
G.varToFac = [[i] for i in range(M)]
```

```
for i in range(N): for var in np.where(H[i]==1)[0]: G.varToFac[var].append(i+M)
```

```
G.factor = [] G.M = M G.N = N G.p = p
```

```
for i in range(M):
```

```
if yhat[i] == 1: vals = [p, 1.-p] if yhat[i] == 0: vals = [1.-p, p] G.factor.append(Factor(None, [i], [2], np.array(vals))) for i in range(N): scope = np.where(H[i]==1)[0]
```

```
val = np.zeros([2]*scope.size) for index, value in np.ndenumerate(val): s = 0 for i in index: s += i if sval[index] == 1.
```

```
G.factor.append(Factor(None, scope, [2]*scope.size, (val)))
```

```
return G
```

XXX

Test cases: $y_{test1} = [1, 1, 1, 1, 1]$ $y_{test2} = [1, 0, 1, 0, 1]$ $y_{test3} = [0, 0, 0, 0, 0]$

If you include the error rate factors: 0, 0, 0.77378094 If you don't: 0, 0, 1

c

All values essentially 0, as expected. Thus, the code word is properly decoded.

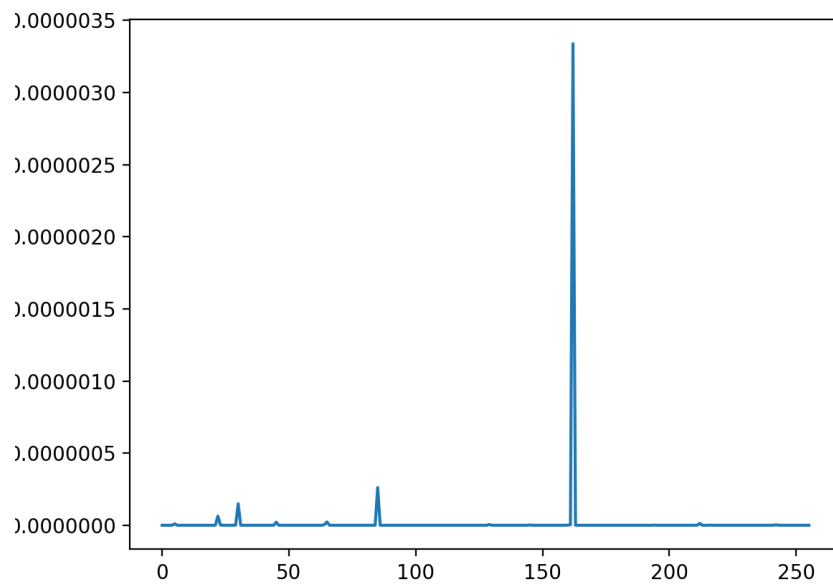


Figure 1: Plot with the i th bit on the X axis and the probability of it being 1 on the y axis.

d

It's moderately reliable. For a 0.06 error rate, all but 1 message converged to the correct message after only 50 iterations.

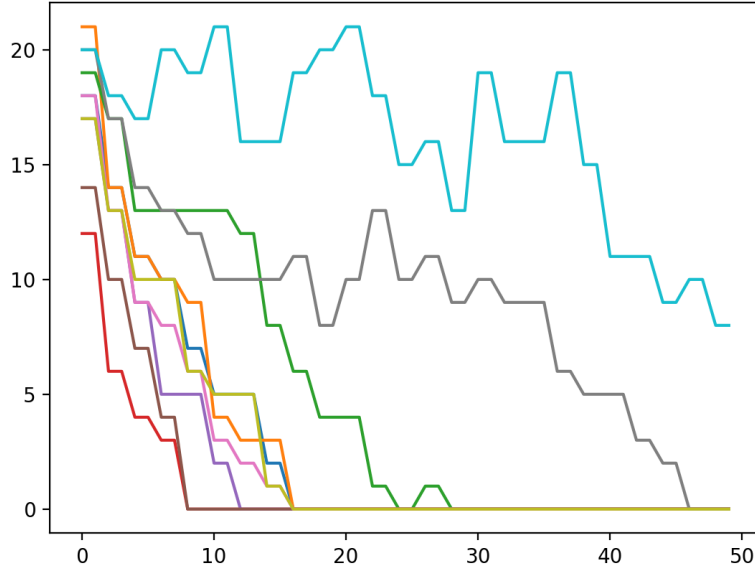


Figure 2: Error rate of 0.06. X axis is the number of iterations. Y axis is the hamming distance between decoded message and real message

e

As the error rate increases, it gets harder and harder for loopy BP to correctly converge onto the correct message. Only 5/10 messages converged to a 0 hamming distance for error rate of 0.08. For the error rate of .1, many didn't even come close and didn't show much sign of going down.

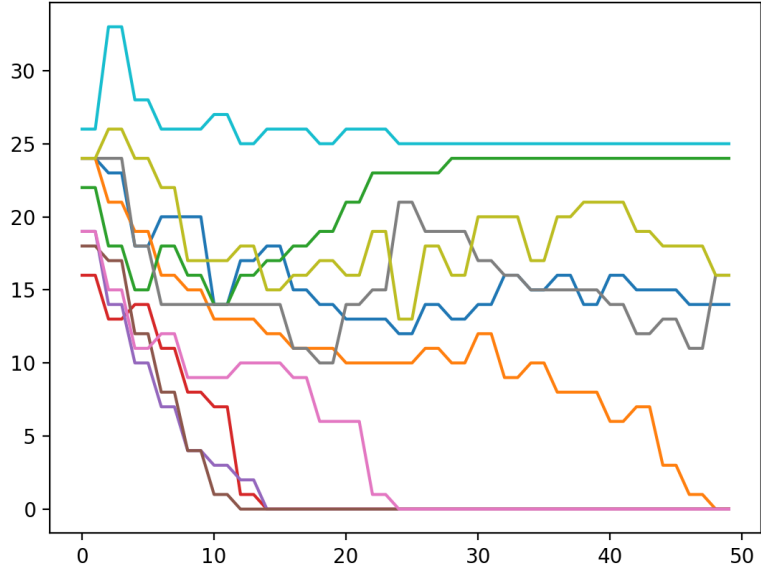


Figure 3: Error rate of 0.08. X axis is the number of iterations. Y axis is the hamming distance between decoded message and real message




Figure 4: Error rate of 0.1. X axis is the number of iterations. Y axis is the hamming distance between decoded message and real message

f

Pretty much decodes picture after 30 iterations. At iteration 0, the picture is non existent since all the incoming messages to variables are initialized uniformly.

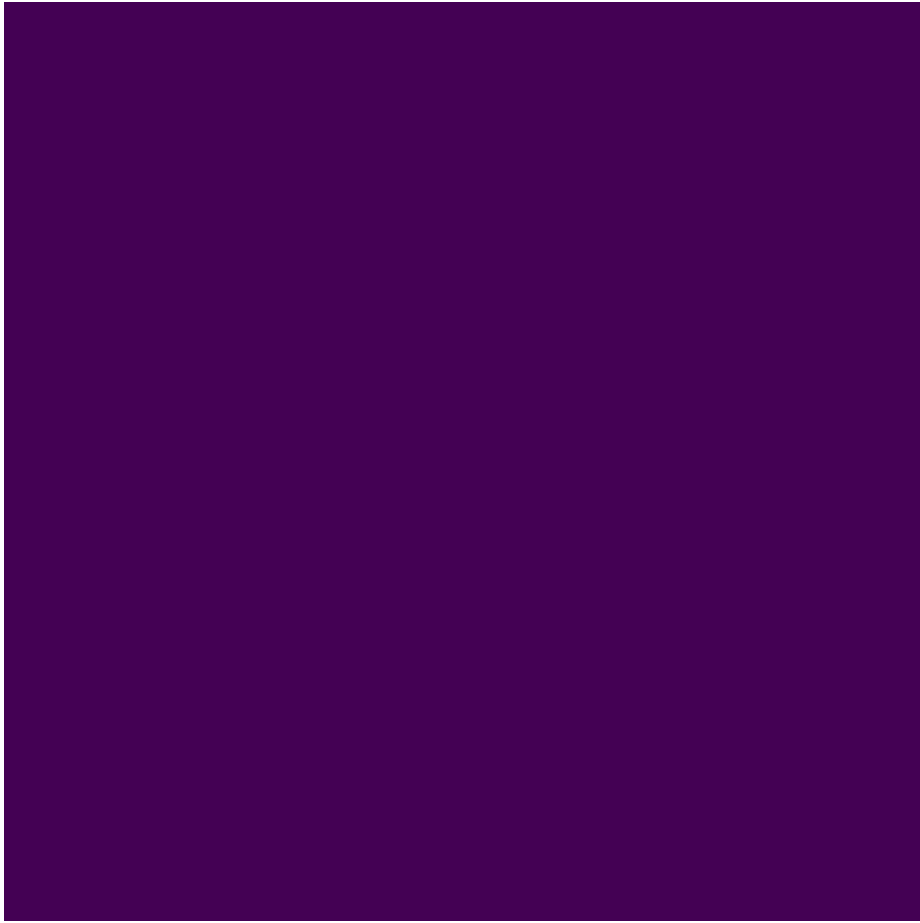


Figure 5: With 0.06 error rate and 0 iteration




Figure 6: With 0.06 error rate and 1 iteration

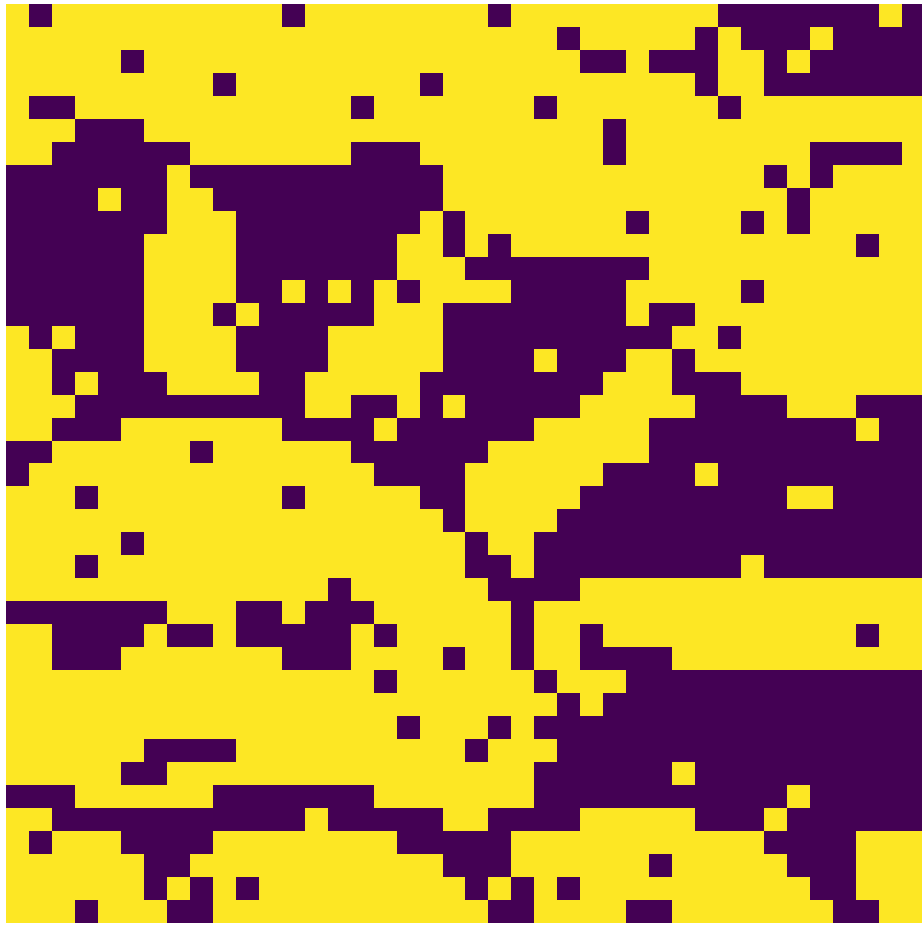


Figure 7: With 0.06 error rate and 2 iteration




Figure 8: With 0.06 error rate and 3 iteration




Figure 9: With 0.06 error rate and 5 iteration

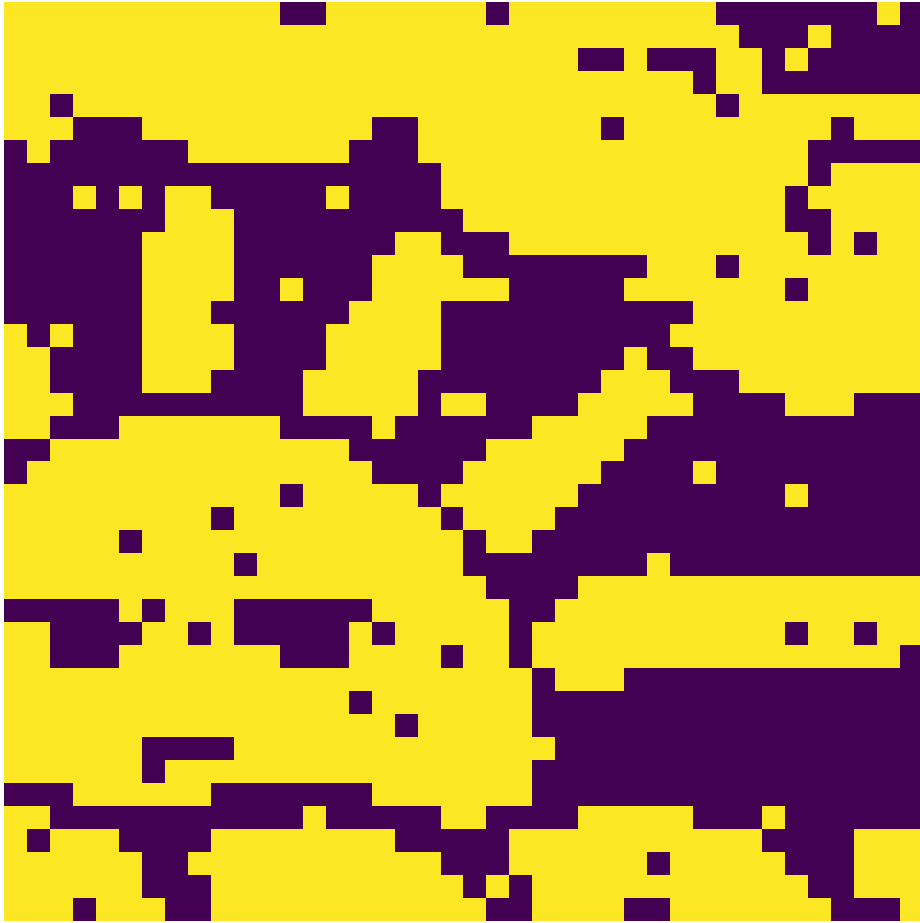


Figure 10: With 0.06 error rate and 10 iteration




Figure 11: With 0.06 error rate and 20 iteration



Figure 12: With 0.06 error rate and 30 iteration

g

At iteration 0, the picture is non existent since all the incoming messages to variables are initialized uniformly. Naturally, since there is a higher error rate, the picture looks more blurry at first, but it doesn't really correct to perfection unlike the error rate of 0.06. At some point, the parity checks get so muffled that you can no longer make out what the picture was supposed to be.

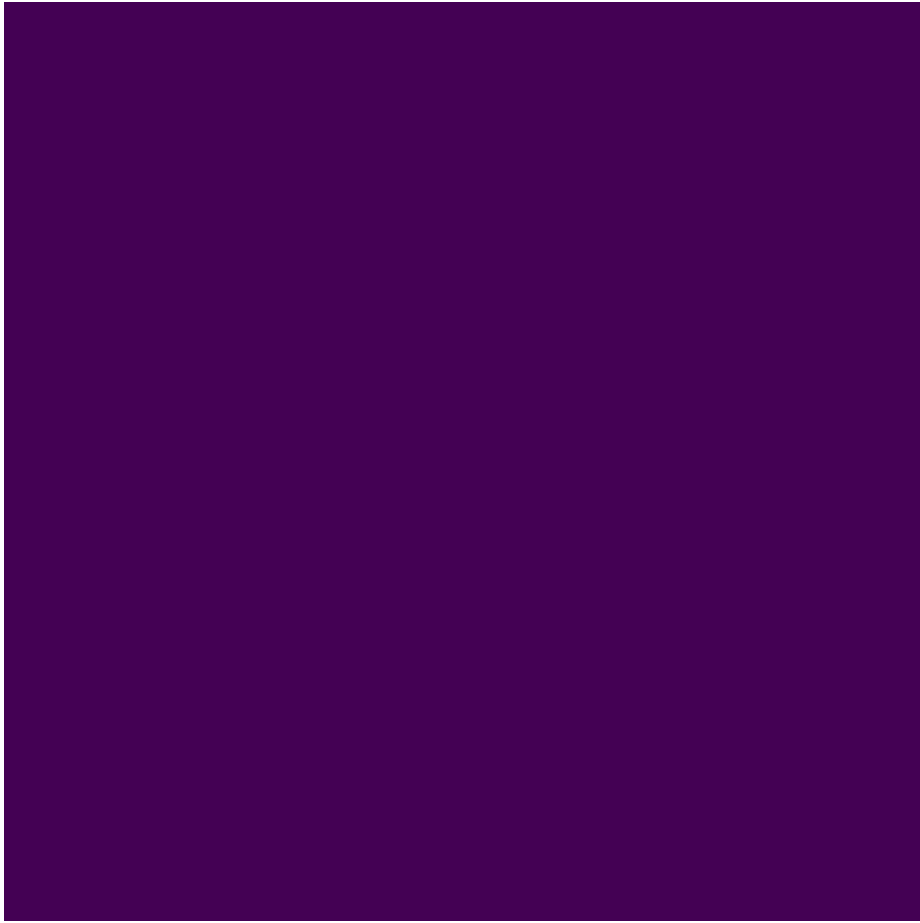


Figure 13: With 0.1 error rate and 0 iteration




Figure 14: With 0.1 error rate and 1 iteration

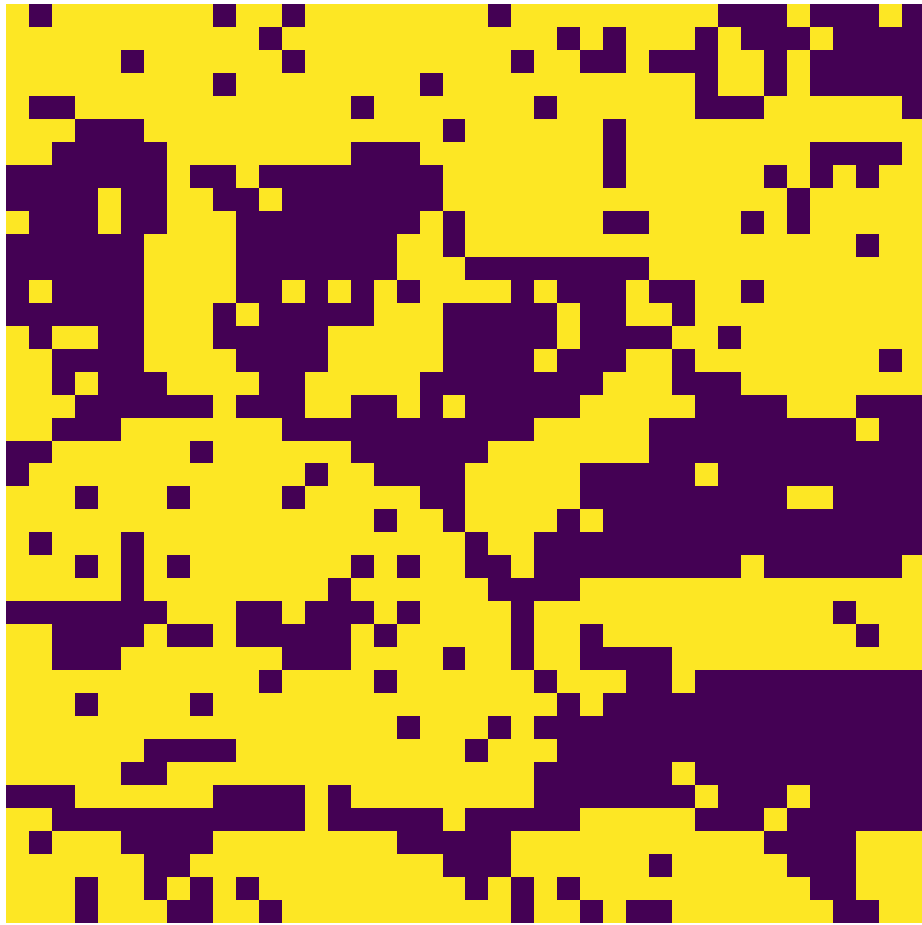


Figure 15: With 0.1 error rate and 2 iteration




Figure 16: With 0.1 error rate and 3 iteration

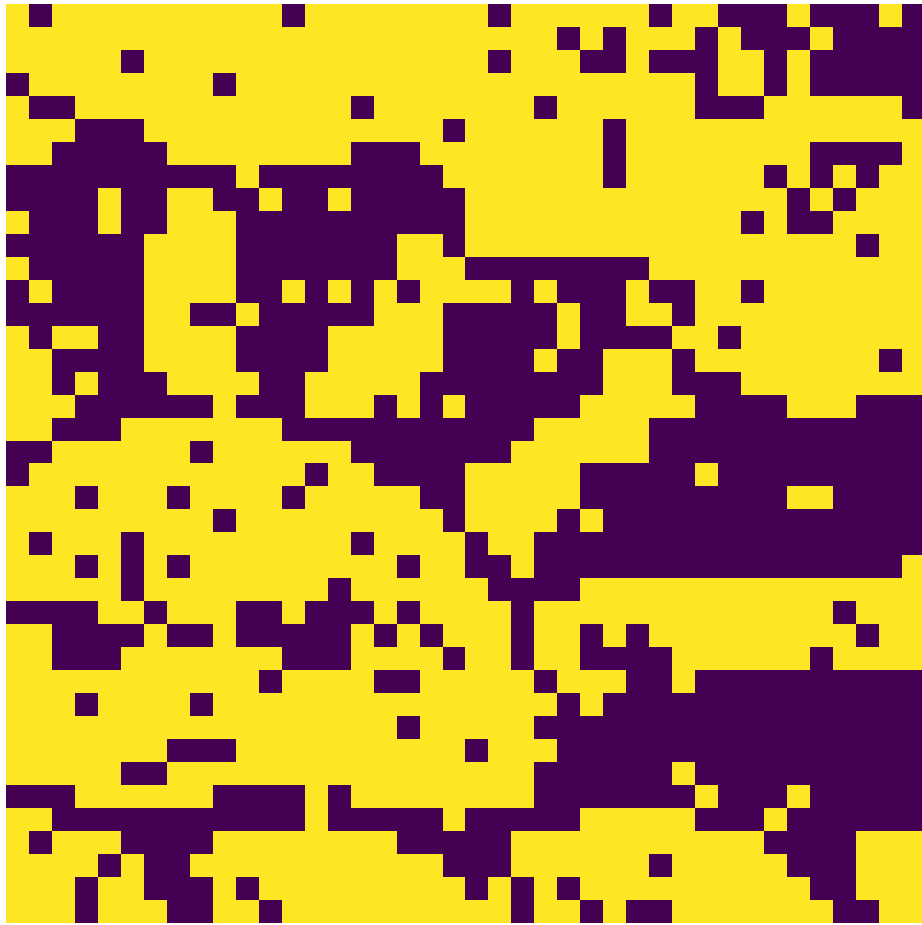


Figure 17: With 0.1 error rate and 5 iteration

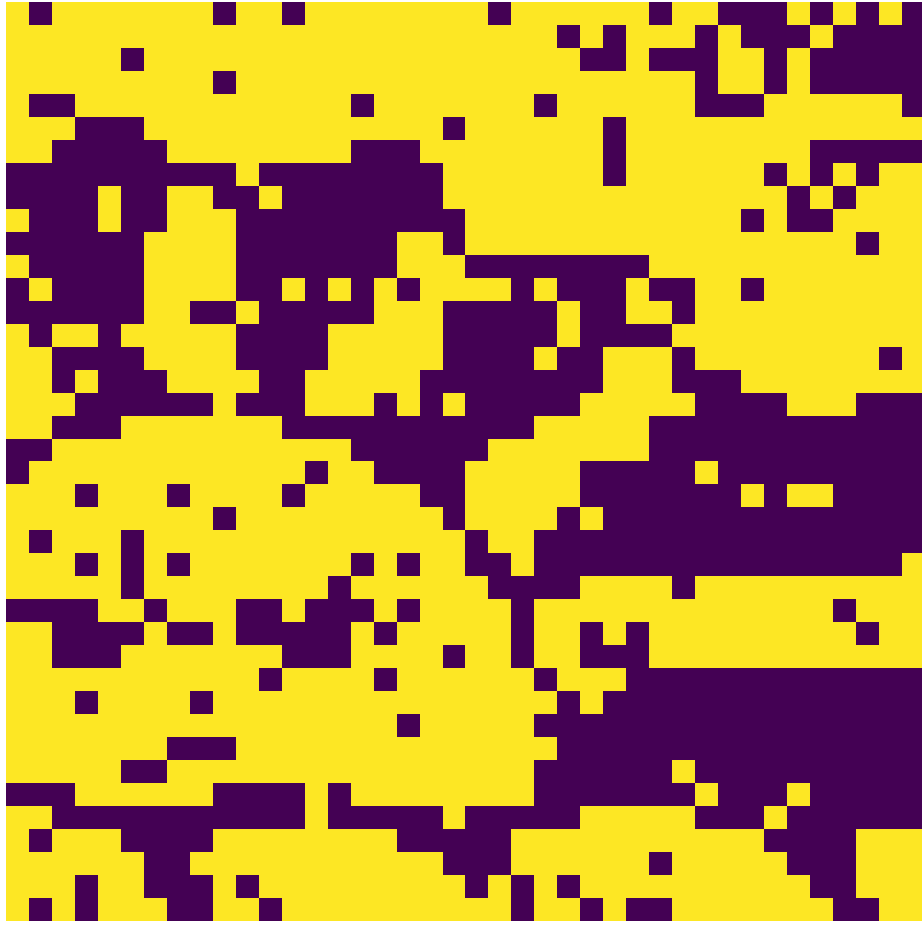


Figure 18: With 0.1 error rate and 10 iteration




Figure 19: With 0.1 error rate and 20 iteration

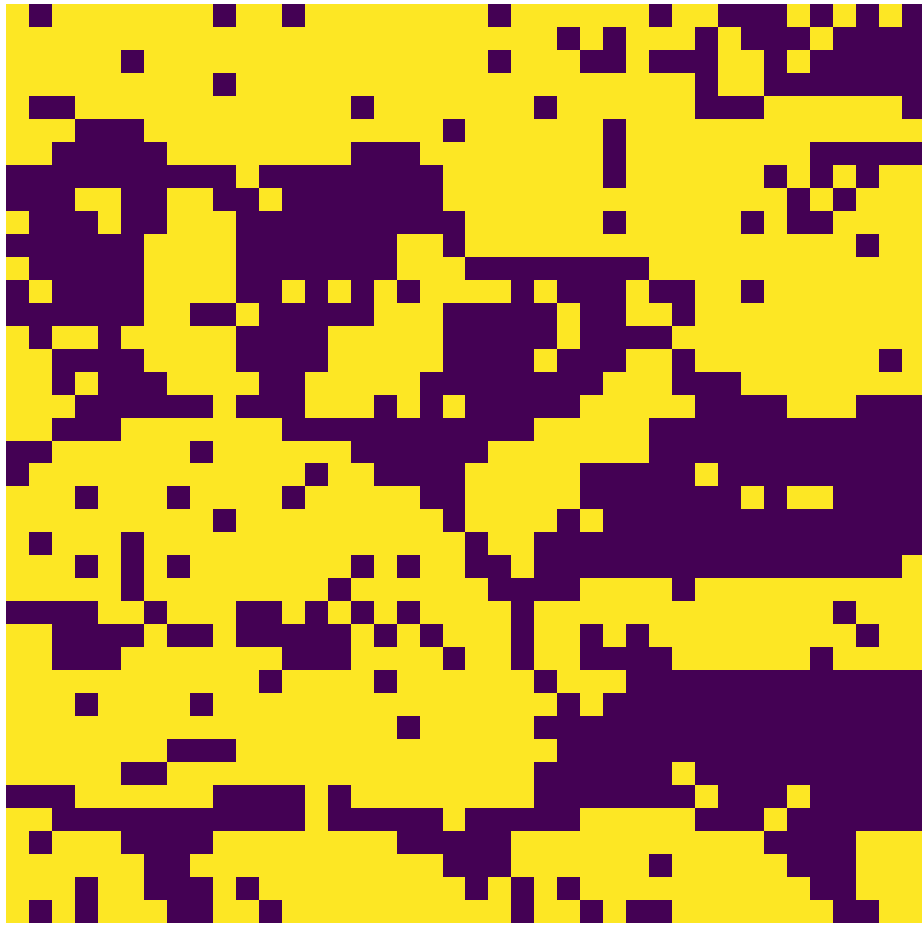


Figure 20: With 0.1 error rate and 30 iteration